

# Memory Simulator

Design Document

## Contents

<b>1 Overview</b>	<b>3</b>
<b>2 System Architecture</b>	<b>3</b>
2.1 Module Structure . . . . .	3
2.2 File Organization . . . . .	4
2.3 Data Flow . . . . .	4
<b>3 Memory Layout and Assumptions</b>	<b>4</b>
3.1 Memory Model . . . . .	4
3.2 Memory Layout Visualization . . . . .	5
3.3 Design Assumptions . . . . .	5
3.4 Why Doubly-Linked List? . . . . .	5
<b>4 Allocation Strategy Implementations</b>	<b>5</b>
4.1 First Fit . . . . .	6
4.2 Best Fit . . . . .	6
4.3 Worst Fit . . . . .	7
4.4 Block Splitting . . . . .	8
4.5 Memory Coalescing . . . . .	8
4.6 Fragmentation Metrics . . . . .	9
<b>5 Cache Hierarchy and Replacement Policy</b>	<b>9</b>
5.1 Cache Architecture . . . . .	9
5.2 Cache Line Structure . . . . .	10
5.3 Address Decomposition . . . . .	10
5.4 Set-Associative Organization . . . . .	11
5.5 Replacement Policies . . . . .	11
5.5.1 FIFO (First-In-First-Out) . . . . .	11
5.5.2 LRU (Least Recently Used) . . . . .	12
5.6 Write Policy Implementation . . . . .	12
5.7 Cache Statistics . . . . .	14
<b>6 Address Translation Flow</b>	<b>14</b>
6.1 Cache Access Flow . . . . .	15
6.2 Multi-Level Access Example . . . . .	16
<b>7 Limitations and Simplifications</b>	<b>16</b>
7.1 Memory Allocator Limitations . . . . .	16
7.2 Cache Simulator Limitations . . . . .	16
7.3 What We Simulate vs. Don't Simulate . . . . .	17
<b>8 Future Extensions</b>	<b>17</b>
8.1 Buddy System Allocator . . . . .	17
8.2 Virtual Memory System . . . . .	18
8.3 Page Replacement Algorithms . . . . .	18
8.4 TLB Simulation . . . . .	18

## 1 Overview

This Memory Management Simulator is an educational tool that demonstrates core operating system concepts related to memory management. It provides hands-on simulation of:

- **Dynamic Memory Allocation:** First Fit, Best Fit, and Worst Fit strategies
- **Memory Fragmentation:** External fragmentation tracking and visualization
- **Cache Simulation:** Multi-level (L1/L2) cache with configurable parameters
- **Cache Replacement Policies:** FIFO and LRU implementations
- **Write Policies:** Write-back with write-allocate, dirty bit tracking

The simulator is implemented in C++17 and provides an interactive command-line interface for experimenting with different configurations.

## 2 System Architecture

### 2.1 Module Structure

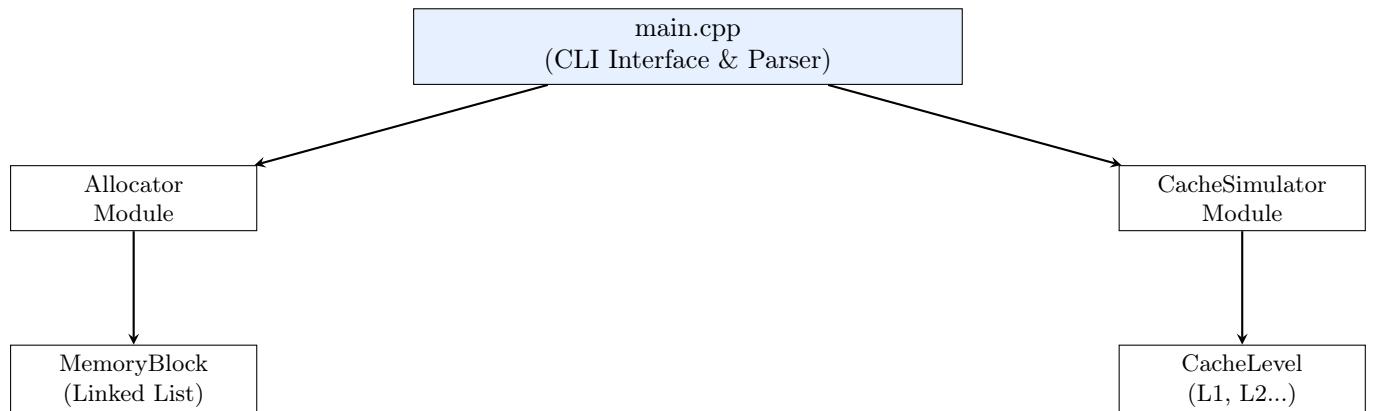
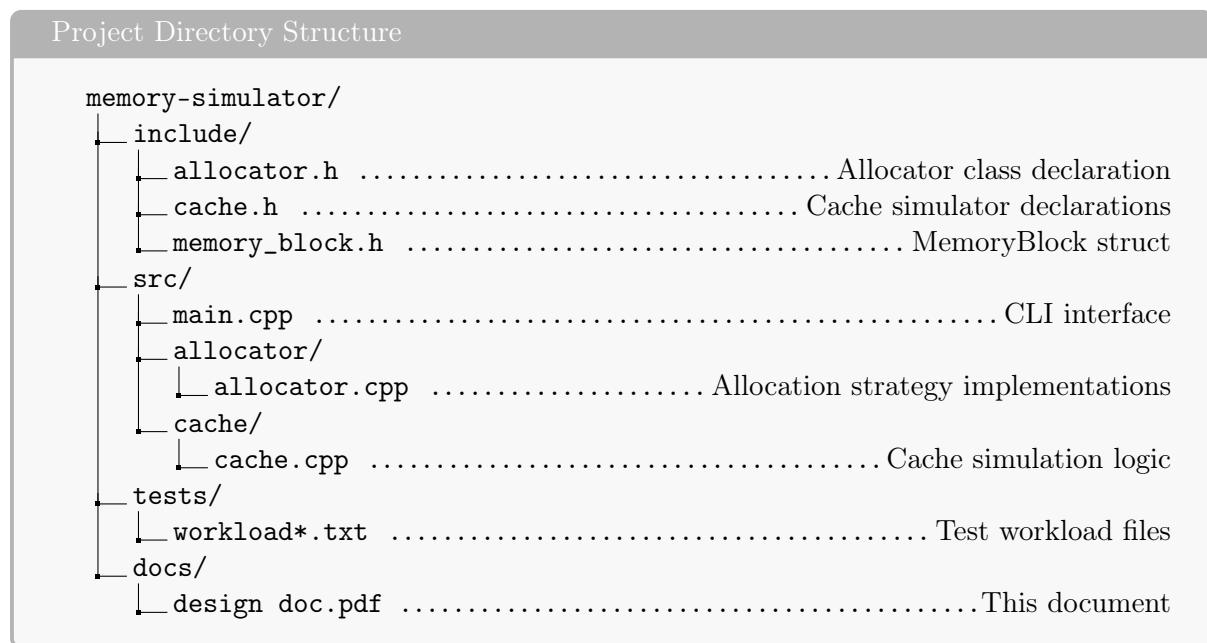
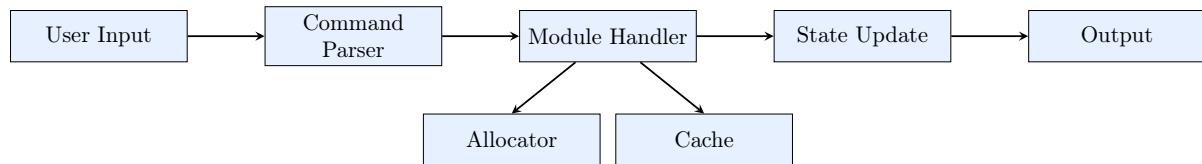


Figure 1: System Architecture Module Structure

## 2.2 File Organization



## 2.3 Data Flow



Workload files can be piped: ./build/memsim < workload.txt

Figure 2: Data Flow Diagram

## 3 Memory Layout and Assumptions

### 3.1 Memory Model

The simulator models physical memory as a contiguous block managed through a **doubly-linked list** of memory blocks:

```

1 struct MemoryBlock {
2     size_t address;           // Starting byte address
3     size_t size;              // Block size in bytes
4     bool is_free;             // Allocation status
5     int block_id;             // Unique ID (-1 if free)
6     MemoryBlock* next;        // Next block pointer
7     MemoryBlock* prev;        // Previous block pointer
8 };

```

Listing 1: MemoryBlock Structure

### 3.2 Memory Layout Visualization

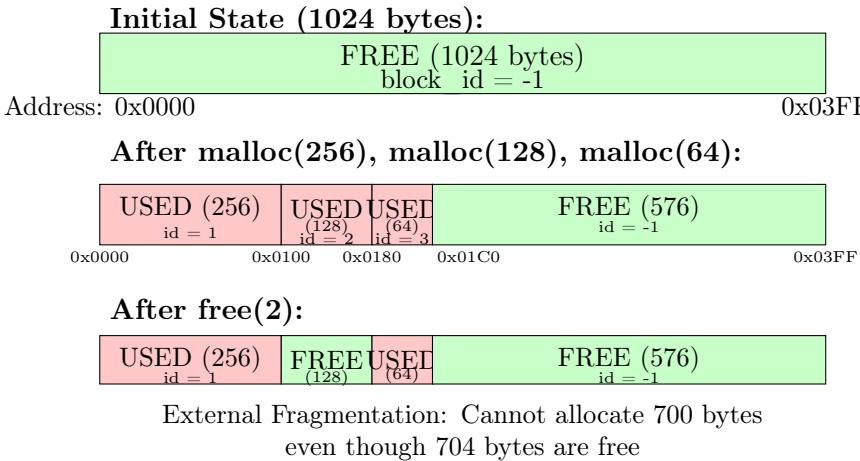


Figure 3: Memory Layout Evolution

### 3.3 Design Assumptions

Assumption	Rationale
Byte-addressable memory	Simplifies allocation math; real systems use word alignment
No alignment requirements	Educational simplicity; real allocators align to 4/8/16 bytes
Instant allocation	No simulation of allocation time; focus on strategy behavior
Unique block IDs	Monotonically increasing; simplifies tracking
No memory protection	No simulation of access rights or segmentation

Table 1: Memory Model Design Assumptions

### 3.4 Why Doubly-Linked List?

Advantages
<ul style="list-style-type: none"> <li>✓ O(1) block splitting: Just adjust pointers</li> <li>✓ O(1) coalescing: Can access both neighbors</li> <li>✓ Dynamic sizing: No pre-defined block count</li> <li>✓ Memory efficient: Only metadata overhead per block</li> </ul>
Trade-offs
<ul style="list-style-type: none"> <li>✗ O(n) search for allocation (traversal required)</li> <li>✗ Pointer overhead per block (16 bytes on 64-bit)</li> </ul>

## 4 Allocation Strategy Implementations

## 4.1 First Fit

**Algorithm:**

```
findBlock(size):
    for each block in list (from head):
        if block.is_free AND block.size >= size:
            return block
    return NULL
```

**Implementation:**

```
1 MemoryBlock* Allocator::firstFit(size_t size) {
2     MemoryBlock* current = head;
3     while (current != nullptr) {
4         if (current->is_free && current->size >= size) {
5             return current;
6         }
7         current = current->next;
8     }
9     return nullptr;
10 }
```

Listing 2: First Fit Implementation

**Characteristics:**

Metric	Value
Time Complexity	O(n) worst, O(1) best
Fragmentation	Tends to fragment start of memory
Best For	General purpose, fast allocation

Table 2: First Fit Characteristics

## 4.2 Best Fit

**Algorithm:**

```
findBlock(size):
    best = NULL
    for each block in list:
        if block.is_free AND block.size >= size:
            if best == NULL OR block.size < best.size:
                best = block
    return best
```

**Implementation:**

```
1 MemoryBlock* Allocator::bestFit(size_t size) {
2     MemoryBlock* best = nullptr;
3     MemoryBlock* current = head;
4     while (current != nullptr) {
5         if (current->is_free && current->size >= size) {
6             if (best == nullptr || current->size < best->size) {
7                 best = current;
8             }
9 }
```

```

9     }
10    current = current->next;
11 }
12 return best;
13 }
```

Listing 3: Best Fit Implementation

**Characteristics:**

Metric	Value
Time Complexity	$O(n)$ always (must scan all)
Fragmentation	Creates many small unusable fragments
Best For	Memory-constrained systems

Table 3: Best Fit Characteristics

### 4.3 Worst Fit

**Algorithm:**

```

findBlock(size):
    worst = NULL
    for each block in list:
        if block.is_free AND block.size >= size:
            if worst == NULL OR block.size > worst.size:
                worst = block
    return worst
```

**Implementation:**

```

1 MemoryBlock* Allocator::worstFit(size_t size) {
2     MemoryBlock* worst = nullptr;
3     MemoryBlock* current = head;
4     while (current != nullptr) {
5         if (current->is_free && current->size >= size) {
6             if (worst == nullptr || current->size > worst->size) {
7                 worst = current;
8             }
9         }
10        current = current->next;
11    }
12    return worst;
13 }
```

Listing 4: Worst Fit Implementation

**Characteristics:**

Metric	Value
Time Complexity	$O(n)$ always
Fragmentation	Leaves larger remaining fragments
Best For	Workloads with similar-sized allocations

Table 4: Worst Fit Characteristics

#### 4.4 Block Splitting

When a block is larger than the requested size, it is split:

**Before split (allocating 100 from 256-byte block):**

FREE (256 bytes)
------------------

**After split:**

USED (100) id = 1	FREE (156) id = -1
----------------------	-----------------------

Figure 4: Block Splitting Example

**Implementation:**

```

1 if (block->size > size) {
2     MemoryBlock* new_free = new MemoryBlock(
3         block->address + size,           // New address
4         block->size - size,            // Remaining size
5         true, -1                     // Free, no ID
6     );
7     // Update linked list pointers...
8     block->size = size;
9 }
```

Listing 5: Block Splitting

#### 4.5 Memory Coalescing

When a block is freed, adjacent free blocks are merged to reduce fragmentation:

**Before free(id=2):**

FREE (100)	USED (200) id=2	FREE (150)	USED (50) id=3
---------------	-----------------------	---------------	----------------------

**After free(id=2) with coalescing:**

FREE (450) (merged block)	USED id=3
------------------------------	--------------

Figure 5: Memory Coalescing Example

**Implementation:**

```

1 void Allocator::coalesce(MemoryBlock* block) {
2     // Merge with next blocks
3     while (block->next && block->next->is_free) {
4         MemoryBlock* next = block->next;
5         block->size += next->size;
6         block->next = next->next;
7         if (next->next) next->next->prev = block;
8         delete next;
9     }
10    // Merge with previous blocks
```

```

11     while (block->prev && block->prev->is_free) {
12         MemoryBlock* prev = block->prev;
13         prev->size += block->size;
14         prev->next = block->next;
15         if (block->next) block->next->prev = prev;
16         delete block;
17         block = prev;
18     }
19 }
```

Listing 6: Memory Coalescing

## 4.6 Fragmentation Metrics

### External Fragmentation Formula:

$$\text{External Fragmentation} = \left(1 - \frac{\text{Largest Free Block}}{\text{Total Free Memory}}\right) \times 100\% \quad (1)$$

**Example:**

**Memory State:**

[USED 100][FREE 50][USED 200][FREE 150][USED 100][FREE 100]

Total Free =  $50 + 150 + 100 = 300$  bytes

Largest Free Block = 150 bytes

External Fragmentation =  $(1 - 150/300) \times 100\% = 50\%$

**Interpretation:** We have 300 bytes free, but cannot allocate anything larger than 150 bytes.

## 5 Cache Hierarchy and Replacement Policy

### 5.1 Cache Architecture

The simulator implements a **multi-level set-associative cache** with configurable parameters:

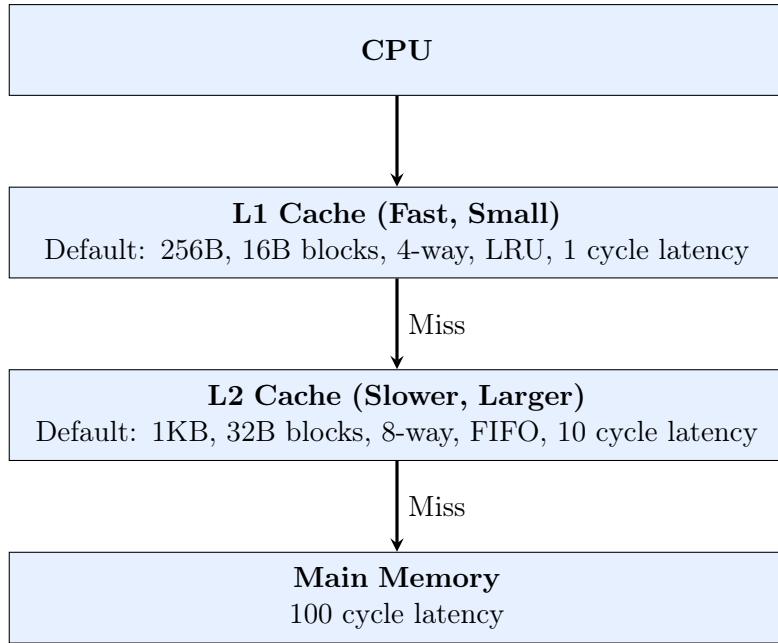


Figure 6: Cache Hierarchy

## 5.2 Cache Line Structure

```

1 struct CacheLine {
2     bool valid;           // Contains valid data?
3     bool dirty;          // Modified since loaded? (for write-back)
4     size_t tag;          // Address tag for matching
5     size_t last_access;  // Access counter for LRU
6 };
  
```

Listing 7: Cache Line Structure

**Valid Bit:** Indicates if the cache line contains meaningful data.

- `valid = false`: Line is empty or invalidated
- `valid = true`: Line contains cached data

**Dirty Bit:** Tracks if data has been written (modified).

- `dirty = false`: Data matches main memory (clean)
- `dirty = true`: Data modified, memory is stale

**Tag:** Upper address bits used to identify which memory block is cached.

## 5.3 Address Decomposition

For a cache with block size  $B$  and  $S$  sets:



Figure 7: Address Bit Decomposition

**Example:** 32-bit address, 16B blocks, 4 sets



Figure 8: 32-bit Address Example

### Implementation:

```

1 size_t CacheLevel::getSetIndex(size_t address) const {
2     size_t block_bits = (size_t)std::log2(block_size); // Offset bits
3     return (address >> block_bits) % num_sets;
4 }
5
6 size_t CacheLevel::getTag(size_t address) const {
7     size_t block_bits = (size_t)std::log2(block_size);
8     size_t set_bits = (size_t)std::log2(num_sets);
9     return address >> (block_bits + set_bits);
10 }
```

Listing 8: Address Decomposition Functions

## 5.4 Set-Associative Organization

### 4-way Set-Associative Cache (4 sets):

Set 0:	Line 0 V D Tag	Line 1 V D Tag	Line 2 V D Tag	Line 3 V D Tag
Set 1:				
Set 2:				
Set 3:				

Figure 9: 4-way Set-Associative Cache Organization

### Cache Mappings:

- **Direct-Mapped:** associativity = 1 (1 line per set)
- **Fully Associative:** num\_sets = 1 (all lines in one set)
- **N-way Set-Assoc:** N lines per set

## 5.5 Replacement Policies

### 5.5.1 FIFO (First-In-First-Out)

Evicts the oldest entry in the set, regardless of access pattern.

Implementation: Queue per set tracking insertion order

Insert A → Insert B → Insert C → Insert D (cache full)  
 Queue: [A, B, C, D]

Insert E (miss, need to evict):  
 - Evict A (front of queue)  
 - Queue: [B, C, D, E]

**Data Structure:**

```

1 std::vector<std::list<size_t>> fifo_queues; // One queue per set
2
3 // On insertion:
4 fifo_queues[set_index].push_back(line_index);
5
6 // On eviction:
7 size_t victim = fifo_queues[set_index].front();
8 fifo_queues[set_index].pop_front();

```

Listing 9: FIFO Implementation

### 5.5.2 LRU (Least Recently Used)

Evicts the line that hasn't been accessed for the longest time.

Implementation: Counter-based tracking

Access A (counter=1) → Access B (2) → Access C (3) → Access A (4)  
 Line A: last\_access = 4  
 Line B: last\_access = 2 ← LRU victim (smallest counter)  
 Line C: last\_access = 3

**Implementation:**

```

1 // On every access:
2 set[line].last_access = ++access_counter;
3
4 // Finding victim:
5 size_t victim = 0;
6 size_t min_access = set[0].last_access;
7 for (size_t i = 1; i < associativity; i++) {
8     if (set[i].last_access < min_access) {
9         min_access = set[i].last_access;
10        victim = i;
11    }
12 }

```

Listing 10: LRU Implementation

## 5.6 Write Policy Implementation

The simulator implements **Write-Back with Write-Allocate**:

Policy	Behavior
<b>Write-Back</b>	Writes go only to cache; memory updated on eviction
<b>Write-Allocate</b>	On write miss, load block into cache then write

Table 5: Write Policy Summary

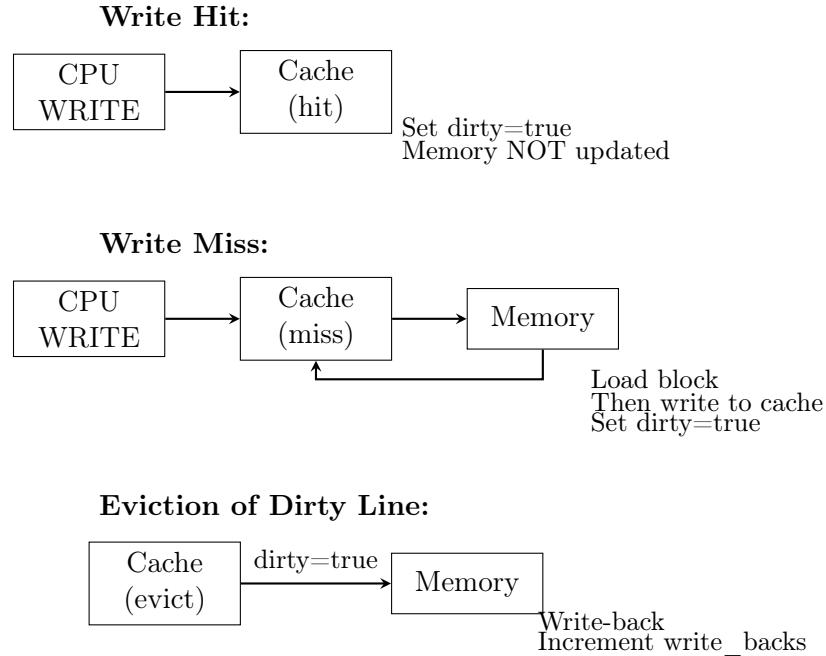


Figure 10: Write Policy Behavior

**Implementation:**

```

1  bool CacheLevel::access(size_t address, bool isWrite) {
2      // ... hit detection ...
3
4      if (hit) {
5          if (isWrite) {
6              set[i].dirty = true; // Mark as modified
7          }
8          return true;
9      }
10
11     // Miss - find victim
12     size_t victim = findVictim(set_index);
13
14     // Write-back if dirty
15     if (set[victim].valid && set[victim].dirty) {
16         stats.write_backs++; // Would write to memory
17     }
18
19     // Load new line
20     set[victim].dirty = isWrite; // Dirty if this is a write
21     // ...
22 }
```

Listing 11: Write Policy Implementation

## 5.7 Cache Statistics

The simulator tracks comprehensive statistics:

```

1 struct CacheStats {
2     size_t hits;           // Cache hits
3     size_t misses;         // Cache misses
4     size_t accesses;       // Total accesses
5     size_t write_backs;    // Dirty evictions
6     size_t total_access_time; // Cumulative latency
7
8     double hitRatio() const {
9         return accesses > 0 ? (double)hits / accesses * 100.0 : 0.0;
10    }
11};

```

Listing 12: Cache Statistics Structure

### Sample Output:

```

==== Cache Statistics ====
L1:
    Accesses:      100
    Hits:          75
    Misses:        25
    Write-backs:   8
    Hit Rate:      75.00%
    Access Time:   100 cycles
L2:
    Accesses:      25
    Hits:          15
    Misses:        10
    Write-backs:   3
    Hit Rate:      60.00%
    Access Time:   250 cycles
-----
Total Access Time: 1350 cycles
Memory Latency:   100 cycles
=====
```

## 6 Address Translation Flow

### 6.1 Cache Access Flow

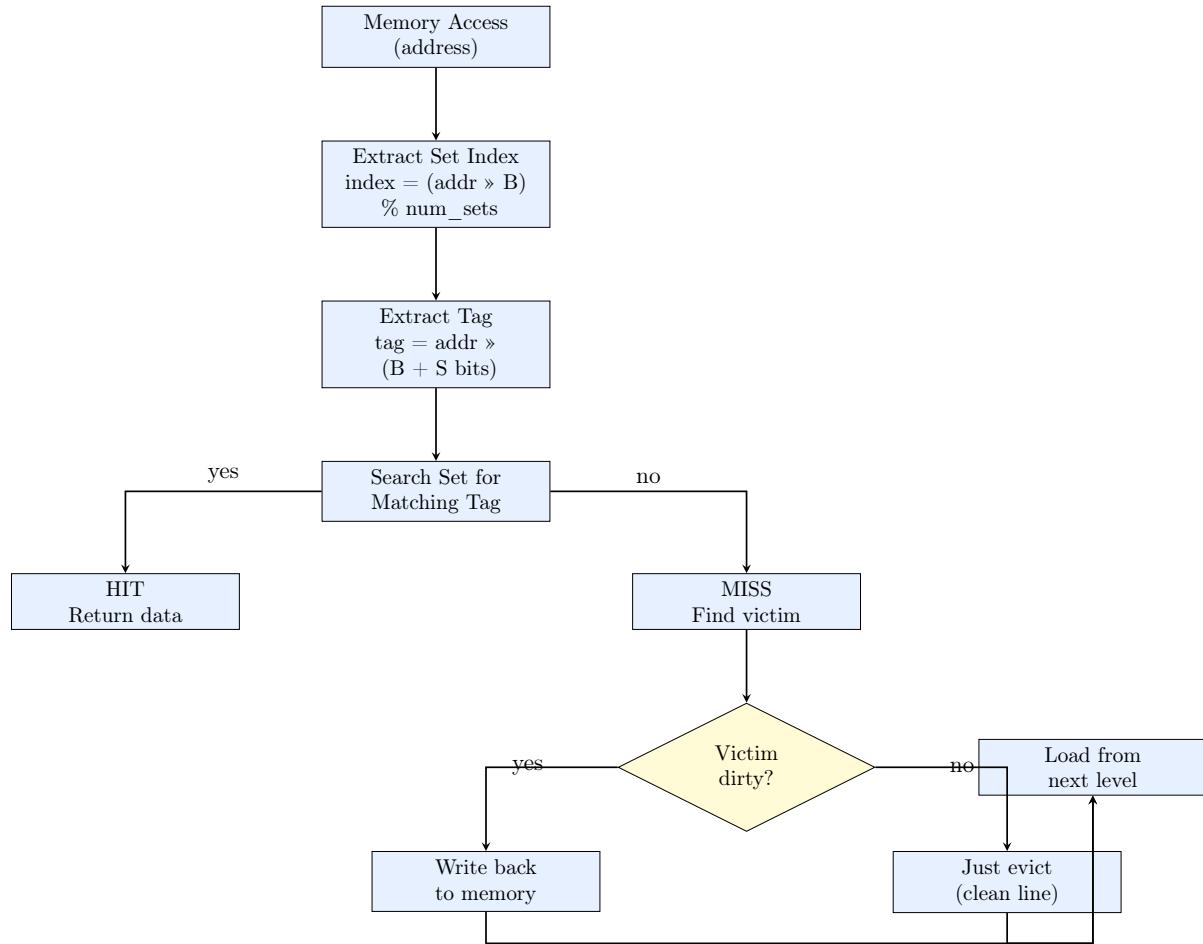


Figure 11: Cache Access Flow Diagram

## 6.2 Multi-Level Access Example

Access address 0x1234:

### Step 1: Check L1

```
Set Index = (0x1234 >> 4) % 4 = 2
Tag = 0x1234 >> 6 = 0x48
→ Search L1 Set 2... MISS
```

### Step 2: Check L2

```
Set Index = (0x1234 >> 5) % 8 = 1
Tag = 0x1234 >> 8 = 0x12
→ Search L2 Set 1... MISS
```

### Step 3: Access Main Memory

→ Fetch block containing 0x1234

### Step 4: Install in caches

→ Load into L2, then L1

**Total Latency:** L1(1) + L2(10) + Memory(100) = **111 cycles**

## 7 Limitations and Simplifications

### 7.1 Memory Allocator Limitations

Limitation	Real-World Behavior
No alignment	Real allocators align to word/cache-line boundaries
No metadata overhead	Real malloc stores size, flags in block headers
Instant operations	Real systems have timing for memory operations
No thread safety	Real allocators handle concurrent access
No memory protection	Real systems have page-level access control

Table 6: Memory Allocator Limitations

### 7.2 Cache Simulator Limitations

Limitation	Real-World Behavior
Exclusive hierarchy	Most caches are inclusive ( $L1 \subseteq L2$ )
No actual data	Only addresses tracked, no data storage
Single CPU	No cache coherence protocols (MESI, etc.)
No prefetching	Real caches speculatively load adjacent blocks
Simplified timing	Real timing includes bus contention, queue delays
Power-of-2 only	Block sizes and set counts must be powers of 2

Table 7: Cache Simulator Limitations

### 7.3 What We Simulate vs. Don't Simulate

✓ Simulated	✗ Not Simulated
Allocation strategies	Virtual memory/paging
Fragmentation tracking	TLB (Translation Lookaside Buffer)
Cache hit/miss behavior	Cache coherence
Replacement policies	Prefetching
Write-back policy	Bus arbitration
Latency accumulation	Actual data storage
Dirty bit tracking	Memory protection

Table 8: Simulation Coverage

## 8 Future Extensions

### 8.1 Buddy System Allocator

Power-of-2 block sizes with efficient splitting/coalescing:

Example: 1024 bytes total

**Request 100 bytes:**

```
1024 → split → 512 + 512
512  → split → 256 + 256
256  → split → 128 + 128
Allocate 128-byte block (smallest power of 2 >= 100)
```

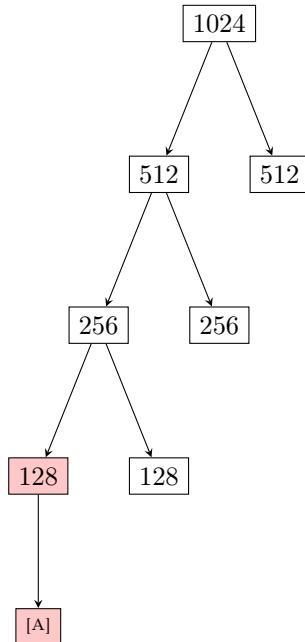


Figure 12: Buddy System Binary Tree

## 8.2 Virtual Memory System

Page table simulation with address translation:

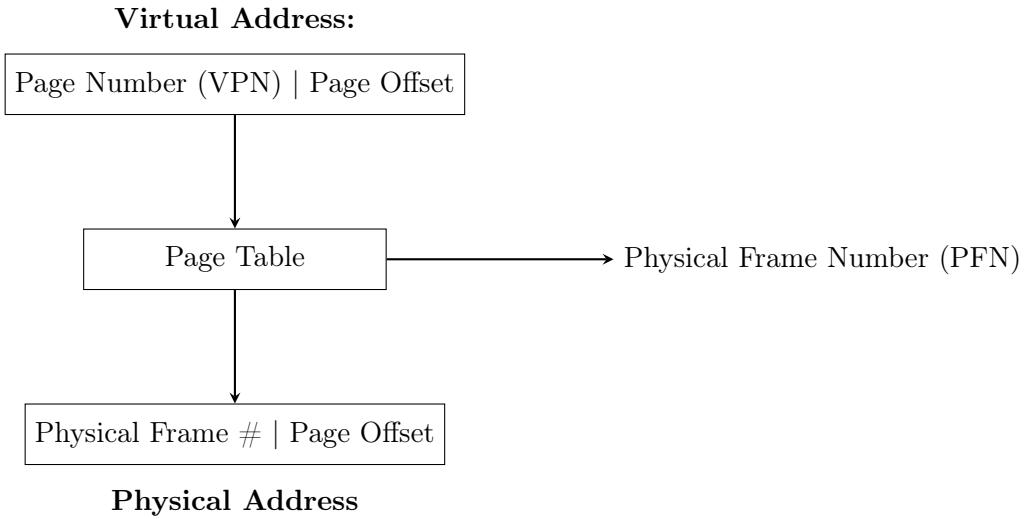


Figure 13: Virtual Memory Address Translation

## 8.3 Page Replacement Algorithms

- **FIFO**: Replace oldest page
- **LRU**: Replace least recently used
- **Clock (Second Chance)**: Circular list with reference bit
- **Optimal**: Replace page used furthest in future (theoretical)

## 8.4 TLB Simulation

Fast address translation cache:

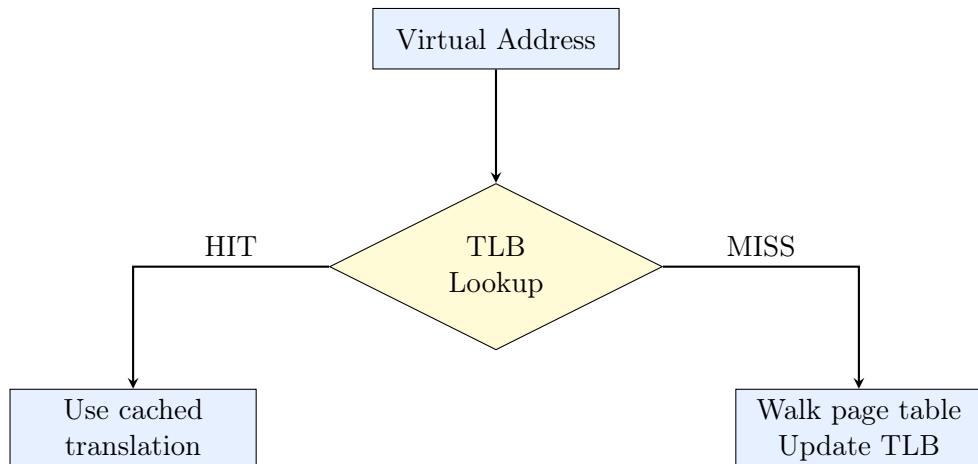


Figure 14: TLB Operation Flow

## Appendix A: Command Reference

Command	Description	Example
init memory <size>	Initialize memory pool	init memory 1024
set allocator <strategy>	Set allocation algorithm	set allocator best_fit
malloc <size>	Allocate memory	malloc 100
free <id>	Free block by ID	free 1
dump memory	Show memory layout	dump memory
stats	Show allocation statistics	stats
init cache	Configure cache (interactive)	init cache
cache read <addr>	Simulate cache read	cache read 0x1000
cache write <addr>	Simulate cache write	cache write 0x1000
cache stats	Show cache statistics	cache stats
cache config	Show cache configuration	cache config
cache reset	Reset cache statistics	cache reset

## Appendix B: Build Instructions

### Build Commands

```
# Release build (optimized)
make

# Debug build (with symbols)
make debug

# Clean build artifacts
make clean

# Run the simulator
./build/memsim

# Run with workload file
./build/memsim < tests/workload1_basic.txt
```

## Appendix C: Test Workloads

The `tests/` directory contains workload files for testing:

File	Purpose
workload1_basic.txt	Basic allocation/deallocation
workload2_strategies.txt	Compare allocation strategies
workload3_cache.txt	Cache hit/miss patterns
workload4_stress.txt	Stress testing
workload5_edge_cases.txt	Boundary conditions
workload6_cache_policies.txt	FIFO vs LRU comparison

Table 10: Test Workload Files